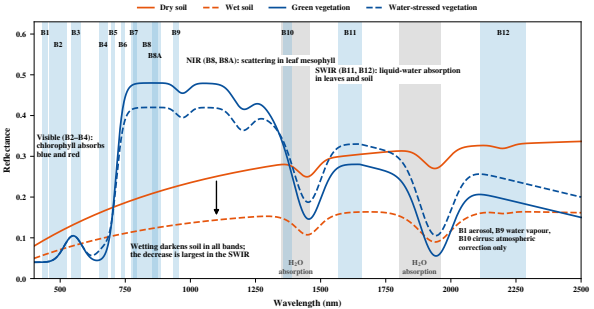
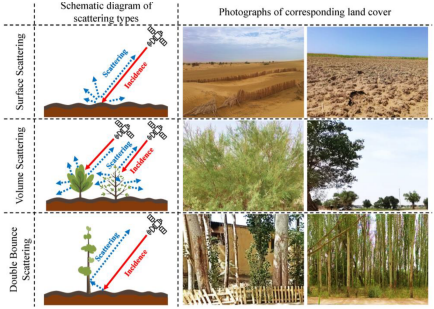
Motivation Different ways of obtaining soil moisture

Source	How SM is obtained	Depth	Main limitation
Microwave (SMAP, SMOS, ASCAT, S1)	invert: σ^0 or $T_B = f(\text{SM}, \text{roughness}, \text{vegetation})$	top $\sim$5 cm; informs part of the root zone (Feldman et al., 2023)	shallow; assumes vegetation and roughness; coarse (9–40 km) except S1
Optical / thermal (S2, Landsat, MODIS)	empirical relations with surface properties: reflectance and indices; LST–vegetation space (e.g. triangle method, TVDI)	skin / surface	indirect, site-specific calibration; clouds; canopy hides the soil
Land-surface models	$d\theta/dt = P - ET - Q - D$, with parameters p	full profile	forcing errors, soil and vegetation parameters, process representation (Koster et al., 2009; Seneviratne et al., 2010)
Reanalysis (ERA5-Land)	LSM driven by reanalysis meteorology	full profile	model assumptions dominate deep layers with few observational constraints (Li et al., 2021; Miguez-Macho and Fan, 2021)
In situ (ISMN)	direct probe measurement	several depths	point scale; sparse, uneven global coverage
Data-driven (ML / DL)	$\text{SM} = f_\theta(\text{inputs})$, learned from in-situ labels; no ad-hoc process assumptions	any depth with labels	needs enough, representative sites; generalisation?; hard to interpret?

Motivation Multimodal remote sensing



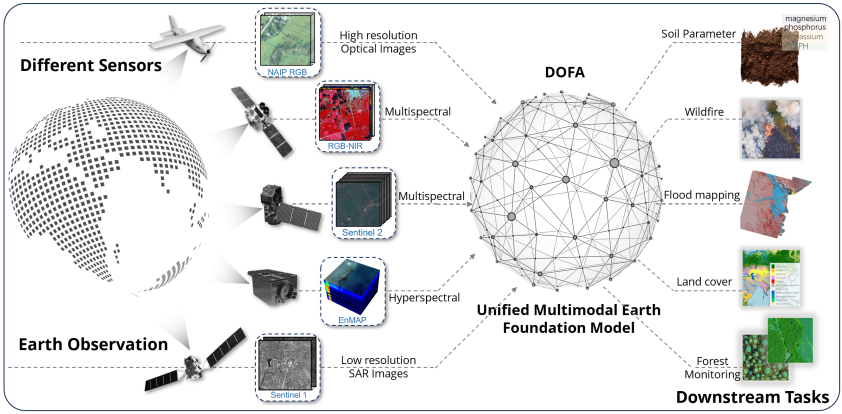
Optical: dry vs wet soil; green vs water-stressed vegetation



SAR scattering types (Nurmemet et al., 2025)

Soil and vegetation water change both optical reflectance and SAR backscatter

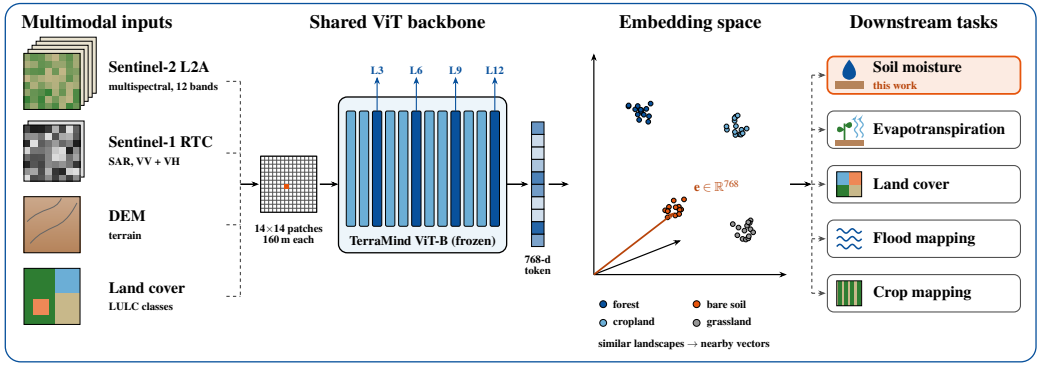
Motivation Geospatial foundation models (GFMs)



DOFA: one model for many sensors and many downstream tasks (Xiong et al., 2024, arXiv:2403.15356)

Once pre-trained on multimodal, unlabelled data, it works for many downstream tasks

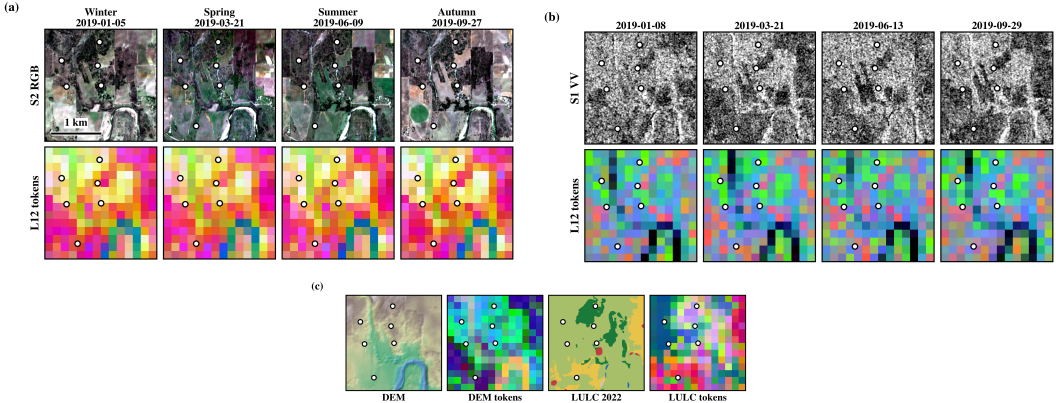
Motivation TerraMind embeddings for soil moisture



S2, S1, DEM and LULC → frozen TerraMind ViT-B → one 768-d embedding per 160 m patch → downstream tasks

The same embedding can serve many tasks; in this case, soil moisture

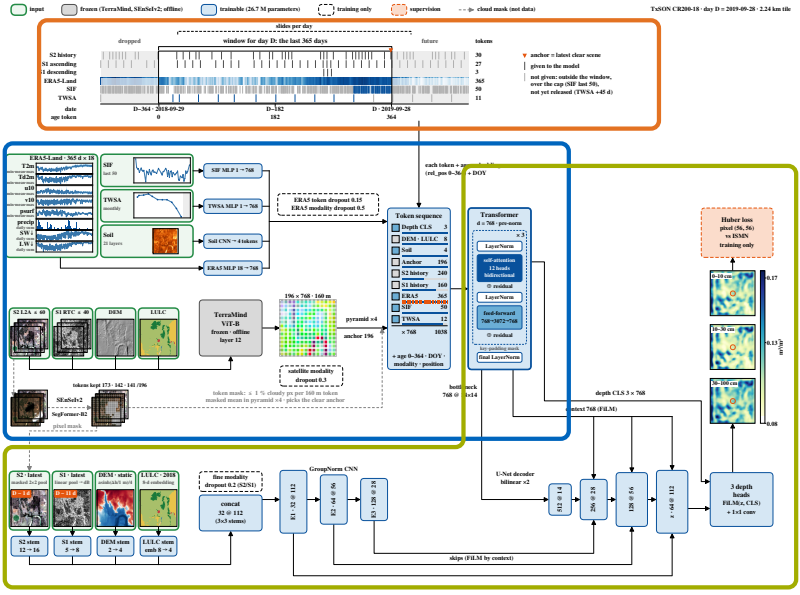
Motivation Example: TerraMind embeddings at one site (TxSON, 2019)



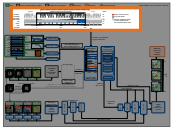
(a) S2 in four seasons, (b) S1 VV, (c) DEM and LULC, each with its TerraMind L12 tokens

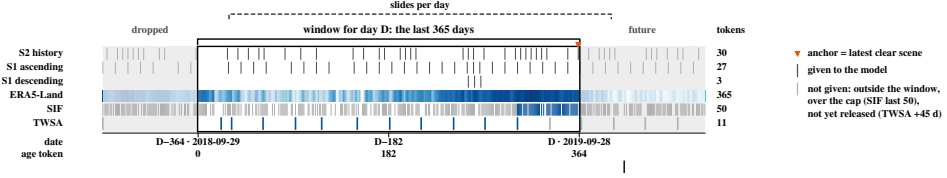
Each 160 m patch → one 768-d vector; colours show its first 3 principal components

Method Model architecture

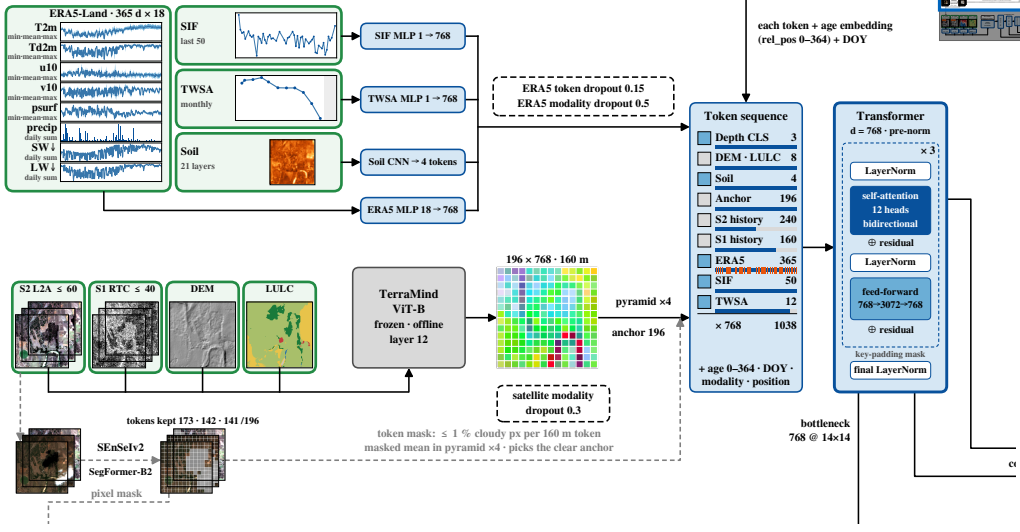


Method Inputs: what the model sees for day D





Method Coarse path: tokens → transformer



Method Fine path: 20 m CNN → U-Net decoder → depth heads

